

case-study-predictive-analysis

November 26, 2024

1 Predictive Modelling

1.1 Linear Regression

1.2 objective

- 1.2.1 A streaming company is focusing on the anime available in their portal and wants to identify the most important factors involved in rating an anime. As a data scientist, you are tasked with analyzing the portal's anime data and identifying the important factors by building a predictive model to predict the rating of an anime.

1.3 Data Dictionary

Each record in the database provides a description of an anime. A detailed data dictionary can be found below.

1. title: title of the anime
2. media Type: format of publication
3. eps: number of episodes (movies are considered 1 episode)
4. duration: duration of an episode in minutes
5. startYr: the year that airing started
6. finishYr: the year that airing finished
7. description: the synopsis of the plot
8. contentWarn: content warning
9. watched: number of users that completed it
10. watching: number of users that are watching it
11. rating: average user rating

12. votes: number of votes that contribute to the rating
13. studio_primary: studios responsible for creation
14. studios_colab: whether there was a collaboration between studios for anime production
15. genre: genre to which the anime belongs

2 Importing Libraries

```
[2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
import statsmodels.api as sm
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

3 Loading Dataset

```
[4]: df = pd.read_csv('/Users/vishal/Desktop/CSV files/anime_rating_dataset.csv')
```

3.0.1 Displaying the first few rows of data set

```
[6]: df.head()
```

```
[6]:
```

		title	mediaType	eps	duration	\
0	Dragon Ball Z Movie 15: Resurrection 'F'		Movie	1	67.0	
1		Kuripuri*Kuripura	Movie	1	5.0	
2		GJ-bu@	TV Special	1	46.0	
3	Nausicaa of the Valley of the Wind		Movie	1	67.0	
4		Miru Tights Special	DVD Special	1	4.0	

	startYr	finishYr	description	\
0	2015	2015	Even the complete obliteration of his physical...	
1	2008	2008		NaN
2	2014	2014	The story is set during the spring vacation im...	
3	1984	1984	One thousand years after the Giant Warriors ca...	
4	2019	2019	Yua is asked to model an illustrator's designs...	

	contentWarn	watched	watching	rating	votes	studio_primary	\
0	No	4649	86	3.979	3100.0	Toei Animation	
1	No	10	0	2.120	10.0	Others	
2	No	1630	16	3.758	1103.0	Others	

3	No	4649	184	4.444	3100.0	Others
4	No	346	8	2.494	234.0	Others

	studios_colab	genre
0	No	Other
1	No	Other
2	No	Other
3	No	Drama
4	No	Other

3.0.2 Checking the shape of the dataset

```
[8]: df.shape
```

```
[8]: (6523, 15)
```

- The dataset contains information (15 attributes) about 6523 anime

3.0.3 Checking the data types of the columns for the dataset

```
[11]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6523 entries, 0 to 6522
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   title                  6523 non-null   object
1   mediaType              6496 non-null   object
2   eps                    6523 non-null   int64
3   duration               6248 non-null   float64
4   startYr                6523 non-null   int64
5   finishYr               6523 non-null   int64
6   description            4114 non-null   object
7   contentWarn            6523 non-null   object
8   watched                6523 non-null   int64
9   watching              6523 non-null   int64
10  rating                 6523 non-null   float64
11  votes                  6496 non-null   float64
12  studio_primary         6523 non-null   object
13  studios_colab          6523 non-null   object
14  genre                  6523 non-null   object
dtypes: float64(3), int64(5), object(7)
memory usage: 764.5+ KB
```

- There are 8 numeric (float and int type) and 7 string (object type) columns in the data

- The target variable is the rating of an anime, which is of float type

4 Statistical summary of the dataset

```
[14]: df.describe()
```

```
[14]:
```

	eps	duration	startYr	finishYr	watched \
count	6523.000000	6248.000000	6523.000000	6523.000000	6523.000000
mean	8.716235	18.396287	2005.241147	2005.575349	1347.948643
std	11.002479	20.949350	12.911035	12.568169	1737.138112
min	1.000000	1.000000	1967.000000	1970.000000	5.000000
25%	1.000000	5.000000	2000.000000	2000.000000	56.000000
50%	1.000000	7.000000	2010.000000	2010.000000	349.000000
75%	12.000000	25.000000	2015.000000	2015.000000	2252.500000
max	34.000000	67.000000	2020.000000	2020.000000	4649.000000

	watching	rating	votes
count	6523.000000	6523.000000	6496.000000
mean	57.445654	2.962553	906.253233
std	76.527405	0.760486	1171.677648
min	0.000000	1.111000	10.000000
25%	2.000000	2.371000	34.000000
50%	13.000000	2.944000	227.500000
75%	98.000000	3.568000	1567.750000
max	199.000000	4.702000	3100.000000

• We can see that the anime ratings vary between 1.1 and 4.7, which suggests that the anime were rated on a scale of 0-5

- The general rating of anime is 2.96
- TV is the most occurring type of media.
- The number of views for the anime in the data varies from 5 to ~5000
- The number of votes received by an anime in the data varies from 10 to 3100

5 Checking for duplicate values

```
[17]: df.duplicated().sum()
```

```
[17]: 0
```

- There are no duplicate values in the data

6 Checking for missing values

```
[20]: df.isnull().sum()
```

```
[20]: title           0
      mediaType      27
      eps           0
      duration      275
      startYr        0
      finishYr       0
      description   2409
      contentWarn    0
      watched       0
      watching      0
      rating        0
      votes         27
      studio_primary 0
      studios_colab  0
      genre         0
      dtype: int64
```

- There are missing values in many columns.

7 Data Preprocessing

```
[23]: df.isnull().sum()
```

```
[23]: title           0
      mediaType      27
      eps           0
      duration      275
      startYr        0
      finishYr       0
      description   2409
      contentWarn    0
      watched       0
      watching      0
      rating        0
      votes         27
      studio_primary 0
      studios_colab  0
      genre         0
      dtype: int64
```

Let's fix the missing values in the data.

- For the variable `mediaType`, we will impute the missing values with ‘Other’ as the exact values for that category are not known

- For the variables `duration` and `votes`, we will impute the missing values in each column with the median grouped by genre and `mediaType` as both the columns have skewed distributions

```
[25]: # we first create a copy of the data to avoid changes to it
df1 = df.copy ()
df1.mediaType.fillna("Other", inplace=True)
df1["duration"] = df1["duration"].fillna(
    value=df1.groupby(["genre", "mediaType"]) ["duration"].transform ("median")
)

df1["votes"] = df1["votes"].fillna(
    value=df1.groupby(["genre", "mediaType"]) ["votes"].transform("median")
)
```

```
[26]: df1.isnull().sum()
```

```
[26]: title                0
mediaType                0
eps                    0
duration                8
startYr                0
finishYr               0
description            2409
contentWarn            0
watched                0
watching               0
rating                 0
votes                  0
studio_primary         0
studios_colab          0
genre                  0
dtype: int64
```

- We will impute the remaining missing values in the `duration` column with the median grouped by genre

```
[28]: df1["duration"] = df1["duration"].fillna(
    value=df1.groupby(["genre"]) ["duration"]. transform("median")
)
```

```
[29]: df1.isnull().sum()
```

```
[29]: title                0
mediaType                0
```

```

eps                0
duration           0
startYr            0
finishYr           0
description        2409
contentWarn        0
watched            0
watching           0
rating             0
votes              0
studio_primary     0
studios_colab      0
genre              0
dtype: int64

```

- All the missing values have been treated.

8 Feature Engineering

- Let's create a new feature `years_running` by taking the difference between `finishYr` and `startyr` columns

- We will drop the original columns once the new feature is created as the year values themselves are not numerical in nature

```

[33]: df1["years_running"] = df1["finishYr"] - df1["startYr"]
      df1.drop(["startYr",
               , "finishYr"], axis=1, inplace=True)
      df1.head()

```

```

[33]:
      title  mediaType  eps  duration  \
0  Dragon Ball Z Movie 15: Resurrection 'F'      Movie      1      67.0
1                Kuripuri*Kuripura      Movie      1       5.0
2                GJ-bu@      TV Special      1      46.0
3  Nausicaa of the Valley of the Wind      Movie      1      67.0
4  Miru Tights Special  DVD Special      1       4.0

      description  contentWarn  watched  \
0  Even the complete obliteration of his physical...      No      4649
1                NaN          No      10
2  The story is set during the spring vacation im...      No      1630
3  One thousand years after the Giant Warriors ca...      No      4649
4  Yua is asked to model an illustrator's designs...      No      346

      watching  rating  votes  studio_primary  studios_colab  genre  \
0          86   3.979  3100.0  Toei Animation          No  Other
1           0   2.120   10.0        Others          No  Other

```

2	16	3.758	1103.0	Others	No	Other
3	184	4.444	3100.0	Others	No	Drama
4	8	2.494	234.0	Others	No	Other

```

years_running
0      0
1      0
2      0
3      0
4      0

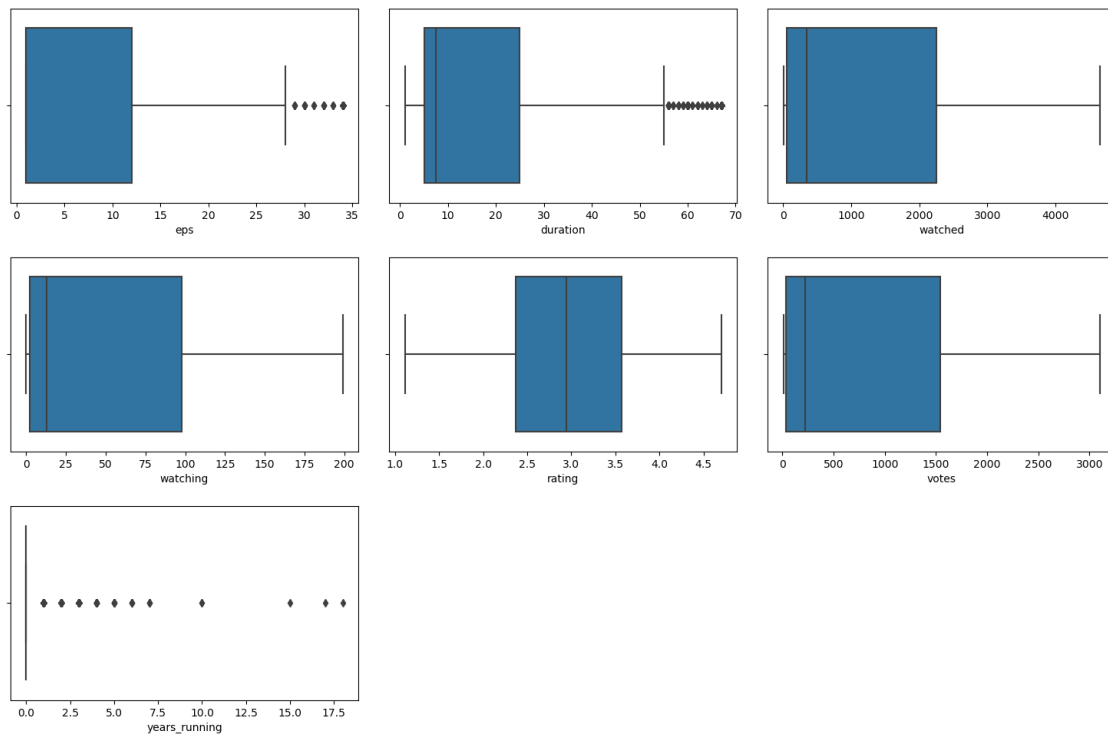
```

9 Outlier detection

```

[35]: # outlier detection using boxplot
num_cols = df1.select_dtypes(include=np.number).columns.tolist()
plt.figure(figsize=(15, 10))
for i, variable in enumerate(num_cols) :
    plt.subplot(3, 3, i + 1)
    sns.boxplot(data=df1, x=variable)
    plt.tight_layout(pad=2)
plt.show()

```



- There are quite a few outliers in the data
- However, we will not treat them as they are proper values
- We will drop the columns title and description as they contain a lot of text and too many unique values, and can be excluded from modeling

```
[37]: df1.drop(["title", "description"], axis=1, inplace=True)
```

```
[38]: df1.head()
```

```
[38]:
```

	mediaType	eps	duration	contentWarn	watched	watching	rating	votes	\
0	Movie	1	67.0	No	4649	86	3.979	3100.0	
1	Movie	1	5.0	No	10	0	2.120	10.0	
2	TV Special	1	46.0	No	1630	16	3.758	1103.0	
3	Movie	1	67.0	No	4649	184	4.444	3100.0	
4	DVD Special	1	4.0	No	346	8	2.494	234.0	

	studio_primary	studios_colab	genre	years_running
0	Toei Animation	No	Other	0
1	Others	No	Other	0
2	Others	No	Other	0
3	Others	No	Drama	0
4	Others	No	Other	0

10 Data Preparation for modeling

- We want to predict the rating of an anime
- Before we proceed to build a model, we'll have to encode categorical features
- We'll split the data into train and test to be able to evaluate the model that we build on the train data
- We will build a Linear Regression model using the train data and then check it's performance

```
[41]: df2=df1.copy()
```

```
[42]: # defining X and y variables
X = df2.drop(["rating"], axis=1)
y = df2["rating"]

# let's add the intercept to data
X = sm.add_constant(X) # Corrected this line

# creating dummy variables
X = pd.get_dummies(
```

```
X,
columns=X.select_dtypes(include=["object", "category"]).columns.tolist(),
drop_first=True
)
X.head()
```

```
[42]:
```

	const	eps	duration	watched	watching	votes	years_running	\
0	1.0	1	67.0	4649	86	3100.0		0
1	1.0	1	5.0	10	0	10.0		0
2	1.0	1	46.0	1630	16	1103.0		0
3	1.0	1	67.0	4649	184	3100.0		0
4	1.0	1	4.0	346	8	234.0		0

	mediaType_Movie	mediaType_Music	Video	mediaType_OVA	...	\
0	True		False	False	...	
1	True		False	False	...	
2	False		False	False	...	
3	True		False	False	...	
4	False		False	False	...	

	studio_primary_Toei Animation	studios_colab_Yes	genre_Adventure	\
0	True	False	False	
1	False	False	False	
2	False	False	False	
3	False	False	False	
4	False	False	False	

	genre_Based on a Manga	genre_Comedy	genre_Drama	genre_Fantasy	\
0	False	False	False	False	
1	False	False	False	False	
2	False	False	False	False	
3	False	False	True	False	
4	False	False	False	False	

	genre_Other	genre_Romance	genre_Sci Fi
0	True	False	False
1	True	False	False
2	True	False	False
3	False	False	False
4	True	False	False

[5 rows x 34 columns]

```
[43]: # converting the input attributes into float type for modeling
x = X.astype(float)
X. head ()
```

```
[43]:
```

	const	eps	duration	watched	watching	votes	years_running	\
0	1.0	1	67.0	4649	86	3100.0		0
1	1.0	1	5.0	10	0	10.0		0
2	1.0	1	46.0	1630	16	1103.0		0
3	1.0	1	67.0	4649	184	3100.0		0
4	1.0	1	4.0	346	8	234.0		0

	mediaType_Movie	mediaType_Music	Video	mediaType_OVA	...	\
0	True		False	False	...	
1	True		False	False	...	
2	False		False	False	...	
3	True		False	False	...	
4	False		False	False	...	

	studio_primary_Toei Animation	studios_colab_Yes	genre_Adventure	\
0	True	False	False	
1	False	False	False	
2	False	False	False	
3	False	False	False	
4	False	False	False	

	genre_Based on a Manga	genre_Comedy	genre_Drama	genre_Fantasy	\
0	False	False	False	False	
1	False	False	False	False	
2	False	False	False	False	
3	False	False	True	False	
4	False	False	False	False	

	genre_Other	genre_Romance	genre_Sci Fi
0	True	False	False
1	True	False	False
2	True	False	False
3	False	False	False
4	True	False	False

[5 rows x 34 columns]

```
[44]: # splitting the data in 70:30 ratio for train to test data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
random_state=1)

print("Number of rows in train data =", X_train.shape[0])
print("Number of rows in test data =", X_test.shape[0])
```

Number of rows in train data = 4566
Number of rows in test data = 1957

11 Model Building - Linear Regression

```
[57]: print(X_train.dtypes)
      print(X_test.dtypes)
```

const	float64
eps	int64
duration	float64
watched	int64
watching	int64
votes	float64
years_running	int64
mediaType_Movie	bool
mediaType_Music Video	bool
mediaType_OVA	bool
mediaType_Other	bool
mediaType_TV	bool
mediaType_TV Special	bool
mediaType_Web	bool
contentWarn_Yes	bool
studio_primary_J.C. Staff	bool
studio_primary_MADHOUSE	bool
studio_primary_OLM	bool
studio_primary_Others	bool
studio_primary_Production I.G	bool
studio_primary_Studio Deen	bool
studio_primary_Studio Pierrot	bool
studio_primary_Sunrise	bool
studio_primary_TMS Entertainment	bool
studio_primary_Toei Animation	bool
studios_colab_Yes	bool
genre_Adventure	bool
genre_Based on a Manga	bool
genre_Comedy	bool
genre_Drama	bool
genre_Fantasy	bool
genre_Other	bool
genre_Romance	bool
genre_Sci Fi	bool
dtype: object	
const	float64
eps	int64
duration	float64
watched	int64
watching	int64
votes	float64
years_running	int64

mediaType_Movie	bool
mediaType_Music Video	bool
mediaType_OVA	bool
mediaType_Other	bool
mediaType_TV	bool
mediaType_TV Special	bool
mediaType_Web	bool
contentWarn_Yes	bool
studio_primary_J.C. Staff	bool
studio_primary_MADHOUSE	bool
studio_primary_OLM	bool
studio_primary_Others	bool
studio_primary_Production I.G	bool
studio_primary_Studio Deen	bool
studio_primary_Studio Pierrot	bool
studio_primary_Sunrise	bool
studio_primary_TMS Entertainment	bool
studio_primary_Toei Animation	bool
studios_colab_Yes	bool
genre_Adventure	bool
genre_Based on a Manga	bool
genre_Comedy	bool
genre_Drama	bool
genre_Fantasy	bool
genre_Other	bool
genre_Romance	bool
genre_Sci Fi	bool
dtype:	object

```
[59]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
import statsmodels.api as sm
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# Convert all boolean columns to integers
X_train = X_train.astype({col: 'int' for col in X_train.
    ↳select_dtypes(include=['bool']).columns})
X_test = X_test.astype({col: 'int' for col in X_test.
    ↳select_dtypes(include=['bool']).columns})

# Ensure all columns are numeric
print(X_train.dtypes)
print(X_test.dtypes)
```

```
# Fit the model again
olsmodel = sm.OLS(y_train, X_train).fit()
print(olsmodel.summary())
```

```
const                                float64
eps                                  int64
duration                            float64
watched                             int64
watching                            int64
votes                                float64
years_running                       int64
mediaType_Movie                     int64
mediaType_Music Video               int64
mediaType_OVA                       int64
mediaType_Other                     int64
mediaType_TV                        int64
mediaType_TV Special                int64
mediaType_Web                       int64
contentWarn_Yes                     int64
studio_primary_J.C. Staff           int64
studio_primary_MADHOUSE             int64
studio_primary_OLM                  int64
studio_primary_Others               int64
studio_primary_Production I.G       int64
studio_primary_Studio Deen          int64
studio_primary_Studio Pierrot       int64
studio_primary_Sunrise              int64
studio_primary_TMS Entertainment    int64
studio_primary_Toei Animation       int64
studios_colab_Yes                   int64
genre_Adventure                     int64
genre_Based on a Manga              int64
genre_Comedy                        int64
genre_Drama                         int64
genre_Fantasy                       int64
genre_Other                         int64
genre_Romance                       int64
genre_Sci Fi                        int64
dtype: object
const                                float64
eps                                  int64
duration                            float64
watched                             int64
watching                            int64
votes                                float64
years_running                       int64
mediaType_Movie                     int64
```

```

mediaType_Music Video      int64
mediaType_OVA              int64
mediaType_Other            int64
mediaType_TV               int64
mediaType_TV Special       int64
mediaType_Web              int64
contentWarn_Yes            int64
studio_primary_J.C. Staff  int64
studio_primary_MADHOUSE    int64
studio_primary_OLM         int64
studio_primary_Others      int64
studio_primary_Production I.G int64
studio_primary_Studio Deen int64
studio_primary_Studio Pierrot int64
studio_primary_Sunrise     int64
studio_primary_TMS Entertainment int64
studio_primary_Toei Animation int64
studios_colab_Yes         int64
genre_Adventure            int64
genre_Based on a Manga     int64
genre_Comedy               int64
genre_Drama                int64
genre_Fantasy              int64
genre_Other                int64
genre_Romance              int64
genre_Sci Fi              int64
dtype: object

```

OLS Regression Results

```

=====
Dep. Variable:          rating    R-squared:          0.722
Model:                  OLS      Adj. R-squared:       0.720
Method:                 Least Squares    F-statistic:      357.4
Date:                  Tue, 26 Nov 2024    Prob (F-statistic): 0.00
Time:                  22:39:51    Log-Likelihood:   -2307.9
No. Observations:      4566    AIC:              4684.
Df Residuals:          4532    BIC:              4902.
Df Model:              33
Covariance Type:       nonrobust
=====

```

```

=====

```

		coef	std err	t	P> t

const		2.7707	0.074	37.657	0.000
2.626	2.915				
eps		0.0193	0.001	17.871	0.000
0.017	0.021				

duration	0.0123	0.000	25.656	0.000
0.011 0.013				
watched	0.0004	2.82e-05	14.213	0.000
0.000 0.000				
watching	0.0037	0.000	21.258	0.000
0.003 0.004				
votes	-0.0003	4.52e-05	-7.460	0.000
-0.000 -0.000				
years_running	-0.0767	0.008	-9.039	0.000
-0.093 -0.060				
mediaType_Movie	-0.2975	0.032	-9.168	0.000
-0.361 -0.234				
mediaType_Music Video	-0.2911	0.030	-9.717	0.000
-0.350 -0.232				
mediaType_OVA	-0.3017	0.030	-10.094	0.000
-0.360 -0.243				
mediaType_Other	-0.2731	0.035	-7.885	0.000
-0.341 -0.205				
mediaType_TV	-0.5301	0.034	-15.800	0.000
-0.596 -0.464				
mediaType_TV Special	-0.1854	0.039	-4.757	0.000
-0.262 -0.109				
mediaType_Web	-0.4087	0.031	-13.263	0.000
-0.469 -0.348				
contentWarn_Yes	-0.1776	0.020	-8.735	0.000
-0.218 -0.138				
studio_primary_J.C. Staff	-0.1578	0.055	-2.850	0.004
-0.266 -0.049				
studio_primary_MADHOUSE	-0.2071	0.057	-3.635	0.000
-0.319 -0.095				
studio_primary_OLM	-0.3696	0.064	-5.785	0.000
-0.495 -0.244				
studio_primary_Others	-0.2482	0.044	-5.603	0.000
-0.335 -0.161				
studio_primary_Production I.G	0.0828	0.058	1.430	0.153
-0.031 0.196				
studio_primary_Studio Deen	-0.1264	0.061	-2.083	0.037
-0.245 -0.007				
studio_primary_Studio Pierrot	-0.2311	0.062	-3.747	0.000
-0.352 -0.110				
studio_primary_Sunrise	-0.0689	0.054	-1.282	0.200
-0.174 0.036				
studio_primary_TMS Entertainment	0.0361	0.057	0.639	0.523
-0.075 0.147				
studio_primary_Toei Animation	-0.1817	0.051	-3.580	0.000
-0.281 -0.082				
studios_colab_Yes	0.0021	0.028	0.072	0.942
-0.054 0.058				

genre_Adventure	-0.1219	0.062	-1.964	0.050
-0.244	-0.000			
genre_Based on a Manga	-0.0056	0.086	-0.065	0.948
-0.175	0.164			
genre_Comedy	-0.2693	0.075	-3.602	0.000
-0.416	-0.123			
genre_Drama	0.2504	0.067	3.740	0.000
0.119	0.382			
genre_Fantasy	0.0621	0.075	0.823	0.411
-0.086	0.210			
genre_Other	-0.0436	0.055	-0.791	0.429
-0.152	0.064			
genre_Romance	0.0026	0.065	0.040	0.968
-0.125	0.130			
genre_Sci Fi	-0.0596	0.068	-0.874	0.382
-0.193	0.074			

Omnibus:	147.177	Durbin-Watson:	1.944
Prob(Omnibus):	0.000	Jarque-Bera (JB):	70.085
Skew:	0.052	Prob(JB):	6.04e-16
Kurtosis:	2.402	Cond. No.	7.64e+04

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 7.64e+04. This might indicate that there are strong multicollinearity or other numerical problems.

12 Interpreting the Regression Results

13 PART - 1

1. Adjusted. R-squared: It reflects the fit of the model.

- Adjusted R-squared values generally range from 0 to 1, where a higher value generally indicates a better fit, assuming certain conditions are met.

- In our case, the value for adj. R-squared is 0.722, which is good.

2. const coefficient: It is the Y-intercept.

- It means that if all the predictor variable coefficients are zero, then the expected output (i.e., Y) would be equal to the const coefficient.

- In our case, the value for const coefficient is 2.7707

3. Coefficient of a predictor variable: It represents the change in the output Y due to a change in the predictor variable (everything else held constant).

- In our case, the coefficient of duration is 0.0123.

14 Model Performance Check

Let's check the performance of the model using different metrics.

- We will be using metric functions defined in sklearn for RMSE, MAE, and R2
- We will define a function to calculate MAPE and adjusted R2
- The mean absolute percentage error (MAPE) measures the accuracy of predictions as a percentage, and can be calculated as the average absolute percent error for each predicted value minus actual values divided by actual values. It works best if there are no extreme values in the data and none of the actual values are 0.
- We will create a function which will print out all the above metrics in one go.

```
[67]: # Function to compute different metrics to check performance of a regression
      ↪ model
def model_performance_regression(model, predictors, target):
    pred = model.predict(predictors)
    r2 = r2_score(target, pred) # Compute R-squared
    rmse = np.sqrt(mean_squared_error(target, pred)) # Compute RMSE
    mae = mean_absolute_error(target, pred) # Compute MAE

    # Creating a dataframe of metrics
    df_perf = pd.DataFrame({
        "RMSE": [rmse],
        "MAE": [mae],
        "R-squared": [r2]
    })
    return df_perf
```

```
[69]: # Checking model performance on train set (seen 70% data)
print("Training Performance\n")
olsmodel_train_perf = model_performance_regression(olsmodel, X_train, y_train)
print(olsmodel_train_perf)
```

Training Performance

	RMSE	MAE	R-squared
0	0.40112	0.330417	0.722387

```
[71]: # Checking model performance on test set (seen 30% data)
print("Test Performance\n")
olsmodel_test_perf = model_performance_regression(olsmodel, X_test, y_test)
print(olsmodel_test_perf)
```

Test Performance

	RMSE	MAE	R-squared
0	0.413096	0.340426	0.703053

Observations

- The training R² is 0.72, so the model is not underfitting
- The train and test RMSE and MAE are comparable, so the model is not overfitting either
- MAE suggests that the model can predict anime ratings within a mean error of 0.34 on the test data

15 2- Checking Linear Regression Assumptions

We will be checking the following Linear Regression assumptions:

1. No Multicollinearity
2. Linearity of variables
3. Independence of error terms
4. Normality of error terms
5. No Heteroscedasticity

16 TEST FOR MULTICOLLINEARITY

- Multicollinearity occurs when predictor variables in a regression model are correlated. This correlation is a problem because predictor variables should be independent. If the correlation between variables is high, it can cause problems when we fit the model and interpret the results. When we have multicollinearity in the linear model, the coefficients that the model suggests are unreliable.
- There are different ways of detecting (or testing) multicollinearity. One such way is by using the Variance Inflation Factor, or VIF.

- **Variance Inflation Factor (VIF):** Variance inflation factors measure the inflation in the variances of the regression parameter estimates due to collinearities that exist among the predictors. It is a measure of how much the variance of the estimated regression coefficient B_k is “inflated” by the existence of correlation among the predictor variables in the model.

- If VIF is 1, then there is no correlation among the k th predictor and the remaining predictor variables, and hence, the variance of B_k is not inflated at all.

- General Rule of thumb:

- If VIF is between 1 and 5, then there is low multicollinearity.

- If VIF is between 5 and 10, we say there is moderate multicollinearity.

- If VIF is exceeding 10, it shows signs of high multicollinearity.

```
[80]: from statsmodels.stats.outliers_influence import variance_inflation_factor
import pandas as pd

def checking_vif(predictors):
    vif = pd.DataFrame()
    vif["feature"] = predictors.columns

    # Calculating VIF for each feature
    vif["VIF"] = [variance_inflation_factor(predictors.values, i) for i in
    ↪range(len(predictors.columns))]

    return vif

# Checking VIF
vif_df = checking_vif(X_train).sort_values('VIF', ascending=False)
print(vif_df)
```

	feature	VIF
0	const	152.488126
5	votes	79.580182
3	watched	68.407244
31	genre_Other	11.591914
18	studio_primary_Others	11.131121
11	mediaType_TV	6.998184
4	watching	4.997799
26	genre_Adventure	4.197895
7	mediaType_Movie	4.062291
24	studio_primary_Toei Animation	3.983035
1	eps	3.907746
32	genre_Romance	3.349247
9	mediaType_OVA	3.049100

29	genre_Drama	2.915488
22	studio_primary_Sunrise	2.881062
2	duration	2.752909
33	genre_Sci Fi	2.698722
23	studio_primary_TMS Entertainment	2.530580
15	studio_primary_J.C. Staff	2.512300
8	mediaType_Music Video	2.409153
16	studio_primary_MADHOUSE	2.359150
19	studio_primary_Production I.G	2.240209
13	mediaType_Web	2.228931
28	genre_Comedy	2.175454
30	genre_Fantasy	2.080246
20	studio_primary_Studio Deen	2.004036
21	studio_primary_Studio Pierrot	2.002522
17	studio_primary_OLM	1.882586
12	mediaType_TV Special	1.756936
10	mediaType_Other	1.735561
27	genre_Based on a Manga	1.687548
6	years_running	1.272615
14	contentWarn_Yes	1.125939
25	studios_colab_Yes	1.043084

- There are multiple columns with very high VIF values, indicating presence of strong multicollinearity
- We will systematically drop numerical columns with $VIF > 5$
- We will ignore the VIF values for dummy variables and the constant (intercept)

17 Remove Multicollinearity

To remove multicollinearity

1. Drop every column one by one that has a VIF score greater than 5.
2. Look at the adjusted R-squared and RMSE of all these models.
3. Drop the variable that makes the least change in adjusted R-squared.
4. Check the VIF scores again.
5. Continue till you get all VIF scores under 5.

Let's define a function that will help us do this.

```
[89]: def treating_multicollinearity(predictors, target, high_vif_columns):
      """
```

Checking the effect of dropping the columns showing high multicollinearity_ on model performance (adj. R-squared and RMSE)

```
predictors: independent variables
target: dependent variable
high_vif_columns: columns having high VIF
"""
```

```
# Empty lists to store adj. R-squared and RMSE values
adj_r2 = []
rmse = []
```

```
# Build OLS models by dropping one of the high VIF columns at a time
# Store the adjusted R-squared and RMSE in the lists defined previously
```

```
for col in high_vif_columns:
    # Defining the new train set
    train = predictors.loc[:, ~predictors.columns.str.startswith(col)]
```

```
    # Create the model
    olsmodel = sm.OLS(target, train).fit()
```

```
    # Adding adj. R-squared and RMSE to the lists
    adj_r2.append(olsmodel.rsquared_adj)
    rmse.append(np.sqrt(olsmodel.mse_resid))
```

```
# Creating a dataframe for the results
temp = pd.DataFrame({
    "col": high_vif_columns,
    "Adj. R-squared after dropping col": adj_r2,
    "RMSE after dropping col": rmse
}).sort_values(by="Adj. R-squared after dropping col", ascending=False)
```

```
temp.reset_index(drop=True, inplace=True)
```

```
return temp
```

```
# Example usage
```

```
col_list = ["watched", "votes"]
res = treating_multicollinearity(X_train, y_train, col_list)
print(res)
```

	col	Adj. R-squared after dropping col	RMSE after dropping col
0	votes	0.716994	0.405042
1	watched	0.707967	0.411451

```
[91]: col_to_drop = "votes"
```

```
# Creating new train and test sets by dropping the specified column
```

```

X_train2 = X_train.loc[:, ~X_train.columns.str.startswith(col_to_drop)]
X_test2 = X_test.loc[:, ~X_test.columns.str.startswith(col_to_drop)]

# Check VIF now
vif = checking_vif(X_train2)
print("VIF after dropping", col_to_drop)
print(vif)

```

VIF after dropping votes

	feature	VIF
0	const	152.249342
1	eps	3.873093
2	duration	2.752030
3	watched	3.235392
4	watching	4.154724
5	years_running	1.272523
6	mediaType_Movie	4.060078
7	mediaType_Music Video	2.408173
8	mediaType_OVA	3.049085
9	mediaType_Other	1.729382
10	mediaType_TV	6.924966
11	mediaType_TV Special	1.756017
12	mediaType_Web	2.227588
13	contentWarn_Yes	1.125939
14	studio_primary_J.C. Staff	2.512299
15	studio_primary_MADHOUSE	2.359013
16	studio_primary_OLM	1.882550
17	studio_primary_Others	11.126123
18	studio_primary_Production I.G	2.240164
19	studio_primary_Studio Deen	2.003836
20	studio_primary_Studio Pierrot	2.002485
21	studio_primary_Sunrise	2.879639
22	studio_primary_TMS Entertainment	2.530283
23	studio_primary_Toei Animation	3.982565
24	studios_colab_Yes	1.042889
25	genre_Adventure	4.197219
26	genre_Based on a Manga	1.687052
27	genre_Comedy	2.175119
28	genre_Drama	2.913752
29	genre_Fantasy	2.079839
30	genre_Other	11.583777
31	genre_Romance	3.348489
32	genre_Sci Fi	2.697440

- We have dealt with multicollinearity in the data
- Let's rebuild the model using the updated set of predictors variables

```
[96]: olsmod1 = sm.OLS(y_train, X_train2).fit()
print(olsmod1.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  rating    R-squared:                  0.719
Model:                          OLS      Adj. R-squared:             0.717
Method:                        Least Squares  F-statistic:                362.4
Date:                          Tue, 26 Nov 2024  Prob (F-statistic):      0.00
Time:                          22:56:23    Log-Likelihood:             -2335.7
No. Observations:              4566      AIC:                       4737.
Df Residuals:                  4533      BIC:                       4950.
Df Model:                      32
Covariance Type:               nonrobust
=====
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                2.7925    0.074    37.755    0.000
2.647    2.937
eps                  0.0200    0.001    18.545    0.000
0.018    0.022
duration             0.0122    0.000    25.374    0.000
0.011    0.013
watched              0.0002    6.17e-06    31.678    0.000
0.000    0.000
watching             0.0031    0.000    19.836    0.000
0.003    0.003
years_running        -0.0761    0.009    -8.922    0.000
-0.093    -0.059
mediaType_Movie      -0.3031    0.033    -9.289    0.000
-0.367    -0.239
mediaType_Music Video -0.2957    0.030    -9.811    0.000
-0.355    -0.237
mediaType_OVA        -0.3012    0.030   -10.017    0.000
-0.360    -0.242
mediaType_Other      -0.2577    0.035    -7.409    0.000
-0.326    -0.189
mediaType_TV         -0.5557    0.034   -16.551    0.000
-0.621    -0.490
mediaType_TV Special -0.1787    0.039    -4.560    0.000
-0.256    -0.102
mediaType_Web        -0.4143    0.031   -13.370    0.000
-0.475    -0.354
contentWarn_Yes      -0.1776    0.020    -8.680    0.000

```


-0.218	-0.137				
studio_primary_J.C. Staff		-0.1579	0.056	-2.835	0.005
-0.267	-0.049				
studio_primary_MADHOUSE		-0.2104	0.057	-3.670	0.000
-0.323	-0.098				
studio_primary_OLM		-0.3675	0.064	-5.717	0.000
-0.494	-0.242				
studio_primary_Others		-0.2552	0.045	-5.728	0.000
-0.343	-0.168				
studio_primary_Production I.G		0.0808	0.058	1.388	0.165
-0.033	0.195				
studio_primary_Studio Deen		-0.1219	0.061	-1.997	0.046
-0.242	-0.002				
studio_primary_Studio Pierrot		-0.2331	0.062	-3.757	0.000
-0.355	-0.111				
studio_primary_Sunrise		-0.0600	0.054	-1.110	0.267
-0.166	0.046				
studio_primary_TMS Entertainment		0.0407	0.057	0.715	0.474
-0.071	0.152				
studio_primary_Toei Animation		-0.1859	0.051	-3.639	0.000
-0.286	-0.086				
studios_colab_Yes		0.0050	0.029	0.173	0.863
-0.051	0.061				
genre_Adventure		-0.1278	0.062	-2.047	0.041
-0.250	-0.005				
genre_Based on a Manga		-0.0166	0.087	-0.191	0.848
-0.187	0.154				
genre_Comedy		-0.2762	0.075	-3.673	0.000
-0.424	-0.129				
genre_Drama		0.2382	0.067	3.538	0.000
0.106	0.370				
genre_Fantasy		0.0542	0.076	0.714	0.475
-0.095	0.203				
genre_Other		-0.0545	0.055	-0.983	0.326
-0.163	0.054				
genre_Romance		-0.0047	0.065	-0.072	0.942
-0.133	0.124				
genre_Sci Fi		-0.0707	0.069	-1.031	0.303
-0.205	0.064				
=====					
Omnibus:	146.326	Durbin-Watson:	1.946		
Prob(Omnibus):	0.000	Jarque-Bera (JB):	69.874		
Skew:	0.053	Prob(JB):	6.71e-16		
Kurtosis:	2.403	Cond. No.	6.35e+04		
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly

specified.

[2] The condition number is large, 6.35e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 6.35e+04. This might indicate that there are strong multicollinearity or other numerical problems.

18 Interpreting the Regression Results:

19 Part 2

4. std err: It reflects the level of accuracy of the coefficients.

- The lower it is, the higher is the level of accuracy.

5. $P > |t|$: It is p-value.

independent variable.

- For each independent feature, there is a null hypothesis and an alternate hypothesis. Here S_{β_i} is the coefficient of the i th feature.
• SH_0 : *Independent feature is not significant* ($S_{\beta_i} = 0$)
• SH_a : *Independent feature is that it is significant* ($S_{\beta_i} \neq 0$)
• $(P > |t|)$ gives the p-value for each independent feature to check that null hypothesis. We are considering 0.05 (5%) as significance level.
• A p-value of less than 0.05 is considered to be statistically significant.
6. Confidence Interval: It represents the range in which our coefficients are likely to fall (with a likelihood of 95%).
Observations
• We can see that adj. R-squared has dropped from 0.720 to 0.717, which shows that the dropped columns did not have much effect on the model
• As there is no multicollinearity, we can look at the p-values of predictor variables to check their significance

20 Dealing with high p-value variables

- Some of the dummy variables in the data have p-value > 0.05 . So, they are not significant and we'll drop them

- But sometimes p-values change after dropping a variable. So, we'll not drop all variables at once

- Instead, we will do the following:

- Build a model, check the p-values of the variables, and drop the column with the highest p-value
- Create a new model without the dropped feature, check the p-values of the variables, and drop the column with the highest p-value
- Repeat the above two steps till there are no columns with p-value > 0.05

Note: The above process can also be done manually by picking one variable at a time that has a high p-value, dropping it, and building a model again. But that might be a little tedious and using a loop will be more efficient.

```
[106]: # Initial list of columns
predictors = X_train2.copy()
cols = predictors.columns.tolist()

# Setting an initial max p-value
max_P_value = 1

while len(cols) > 0:
    # Defining the train set
    X_train_aux = predictors[cols]

    # Fitting the model
    model = sm.OLS(y_train, X_train_aux).fit()

    # Getting the p-values and the maximum p-value
    p_values = model.pvalues
    max_P_value = max(p_values)

    # Name of the variable with maximum p-value
    feature_with_P_max = p_values.idxmax()

    if max_P_value > 0.05:
        cols.remove(feature_with_P_max)
    else:
        break

selected_features = cols
print(selected_features)
```

```
['const', 'eps', 'duration', 'watched', 'watching', 'years_running',
'mediaType_Movie', 'mediaType_Music Video', 'mediaType_OVA', 'mediaType_Other',
'mediaType_TV', 'mediaType_TV Special', 'mediaType_Web', 'contentWarn_Yes',
'studio_primary_J.C. Staff', 'studio_primary_MADHOUSE', 'studio_primary_OLM',
'studio_primary_Others', 'studio_primary_Studio Deen', 'studio_primary_Studio
Pierrot', 'studio_primary_Sunrise', 'studio_primary_Toei Animation',
'genre_Adventure', 'genre_Comedy', 'genre_Drama']
```

```
[108]: # Using selected features to create new train and test sets
```

```
X_train3 = X_train2[selected_features]
```

```
X_test3 = X_test2[selected_features]
```

```
# Fitting the OLS model with the new training set
```

```
olsmod2 = sm.OLS(y_train, X_train3).fit()
```

```
print(olsmod2.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          rating    R-squared:                0.718
Model:                  OLS      Adj. R-squared:           0.717
Method:                 Least Squares    F-statistic:            482.7
Date:                   Tue, 26 Nov 2024    Prob (F-statistic):      0.00
Time:                   23:04:49    Log-Likelihood:         -2340.4
No. Observations:       4566    AIC:                    4731.
Df Residuals:           4541    BIC:                    4891.
Df Model:                24
Covariance Type:        nonrobust
=====
```

```
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                2.7875      0.033    84.367      0.000
2.723      2.852
eps                  0.0201      0.001    18.671      0.000
0.018      0.022
duration             0.0123      0.000    26.100      0.000
0.011      0.013
watched              0.0002    6.16e-06    31.750      0.000
0.000      0.000
watching             0.0031      0.000    19.918      0.000
0.003      0.003
years_running       -0.0762      0.009    -8.944      0.000
-0.093     -0.059
mediaType_Movie     -0.3078      0.032    -9.550      0.000
-0.371     -0.245
mediaType_Music Video -0.2987      0.030    -9.971      0.000
-0.357     -0.240
mediaType_OVA       -0.3016      0.030   -10.092      0.000
-0.360     -0.243
mediaType_Other     -0.2607      0.035    -7.548      0.000
-0.328     -0.193
mediaType_TV        -0.5598      0.033   -16.781      0.000
-0.625     -0.494
```

mediaType_TV Special	-0.1816	0.039	-4.657	0.000
-0.258 -0.105				
mediaType_Web	-0.4164	0.031	-13.523	0.000
-0.477 -0.356				
contentWarn_Yes	-0.1786	0.020	-8.745	0.000
-0.219 -0.139				
studio_primary_J.C. Staff	-0.2017	0.041	-4.865	0.000
-0.283 -0.120				
studio_primary_MADHOUSE	-0.2532	0.043	-5.844	0.000
-0.338 -0.168				
studio_primary_OLM	-0.4121	0.052	-7.941	0.000
-0.514 -0.310				
studio_primary_Others	-0.2993	0.024	-12.445	0.000
-0.346 -0.252				
studio_primary_Studio Deen	-0.1657	0.048	-3.432	0.001
-0.260 -0.071				
studio_primary_Studio Pierrot	-0.2740	0.049	-5.591	0.000
-0.370 -0.178				
studio_primary_Sunrise	-0.1088	0.039	-2.815	0.005
-0.185 -0.033				
studio_primary_Toei Animation	-0.2325	0.034	-6.853	0.000
-0.299 -0.166				
genre_Adventure	-0.0783	0.032	-2.467	0.014
-0.140 -0.016				
genre_Comedy	-0.2271	0.052	-4.374	0.000
-0.329 -0.125				
genre_Drama	0.2881	0.040	7.192	0.000
0.210 0.367				

```

=====
Omnibus:                143.175   Durbin-Watson:                1.946
Prob(Omnibus):          0.000   Jarque-Bera (JB):            68.833
Skew:                   0.052   Prob(JB):                    1.13e-15
Kurtosis:               2.407   Cond. No.                    2.81e+04
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.81e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, $2.81e+04$. This might indicate that there are strong multicollinearity or other numerical problems.

```
[111]: # Checking model performance on train set (seen 70% data)
print("Training Performance\n")
olsmod2_train_perf = model_performance_regression(olsmod2, X_train3, y_train)
print(olsmod2_train_perf)
```

Training Performance

	RMSE	MAE	R-squared
0	0.40399	0.33291	0.7184

```
[113]: # Checking model performance on test set (seen 30% data)
print("Test Performance\n")
olsmod2_test_perf = model_performance_regression(olsmod2, X_test3, y_test)
print(olsmod2_test_perf)
```

Test Performance

	RMSE	MAE	R-squared
0	0.414583	0.341345	0.700912

Observations

- Now no feature has p-value greater than 0.05, so we'll consider the features in `x_train3` as the final set of predictor variables and `olsmod2` as the final model to move forward with
- Now adjusted R-squared is 0.717, i.e., our model is able to explain ~72% of the variance
- The adjusted R-squared in `olsmod1` (where we considered the variables without multicollinearity) was 0.717
- This shows that the variables we dropped were not affecting the model
- RMSE and MAE values are comparable for train and test sets, indicating that the model is not overfitting

Now we'll check the rest of the assumptions on `olsmod2`.

2. Linearity of variables

3. Independence of error terms

4. Normality of error terms

5. No Heteroscedasticity

21 TEST FOR LINEARITY AND INDEPENDENCE

21.0.1 Why the test?

- Linearity describes a straight-line relationship between two variables, predictor variables must have a linear relation with the dependent variable.
- The independence of the error terms (or residuals) is important. If the residuals are not independent, then the confidence intervals of the coefficient estimates will be narrower and make us incorrectly conclude a parameter to be statistically significant.

21.0.2 How to check linearity and independence?

- Make a plot of fitted values vs residuals.
- If they don't follow any pattern, then we say the model is linear and residuals are independent.
- Otherwise, the model is showing signs of non-linearity and residuals are not independent.

21.0.3 How to fix if this assumption is not followed?

- We can try to transform the variables and make the relationships linear.


```
[119]: # Create a dataframe with actual, fitted, and residual values
df_pred = pd.DataFrame()
df_pred["Actual Values"] = y_train # actual values
df_pred["Fitted Values"] = olsmod2.fittedvalues # predicted values
df_pred["Residuals"] = olsmod2.resid # residuals
print(df_pred.head())
```

	Actual Values	Fitted Values	Residuals
5432	2.872	2.795321	0.076679
5326	2.766	2.275887	0.490113
1021	4.049	4.446845	-0.397845
836	3.153	3.176604	-0.023604
1396	2.167	2.265921	-0.098921

```
[121]: import seaborn as sns
import matplotlib.pyplot as plt

# Let's plot the fitted values vs residuals
sns.residplot(data=df_pred, x="Fitted Values", y="Residuals", color="purple",
              lowess=True)
plt.xlabel("Fitted Values")
plt.ylabel("Residuals")
```

```
plt.title("Fitted vs Residual plot")
plt.show()
```



- The scatter plot shows the distribution of residuals (errors) vs fitted values (predicted values).
- If there exist any pattern in this plot, we consider it as signs of non-linearity in the data and a pattern means that the model doesn't capture non-linear effects.
- We see no pattern in the plot above. Hence, the assumptions of linearity and independence are satisfied.

22 Test For Normality

Why the test?

- Error terms, or residuals, should be normally distributed. If the error terms are not normally distributed, confidence intervals of the coefficient estimates may become too wide or narrow. Once confidence interval becomes unstable, it leads to difficulty in

estimating coefficients based on minimization of least squares. Non-normality suggests that there are a few unusual data points that must be studied closely to make a better model.

How to check normality?

- The shape of the histogram of residuals can give an initial idea about the normality.
- It can also be checked via a Q-Q plot of residuals. If the residuals follow a normal distribution, they will make a straight line plot, otherwise not.
- Other tests to check for normality includes the Shapiro-Wilk test.
- Null hypothesis: Residuals are normally distributed
- Alternate hypothesis: Residuals are not normally distributed

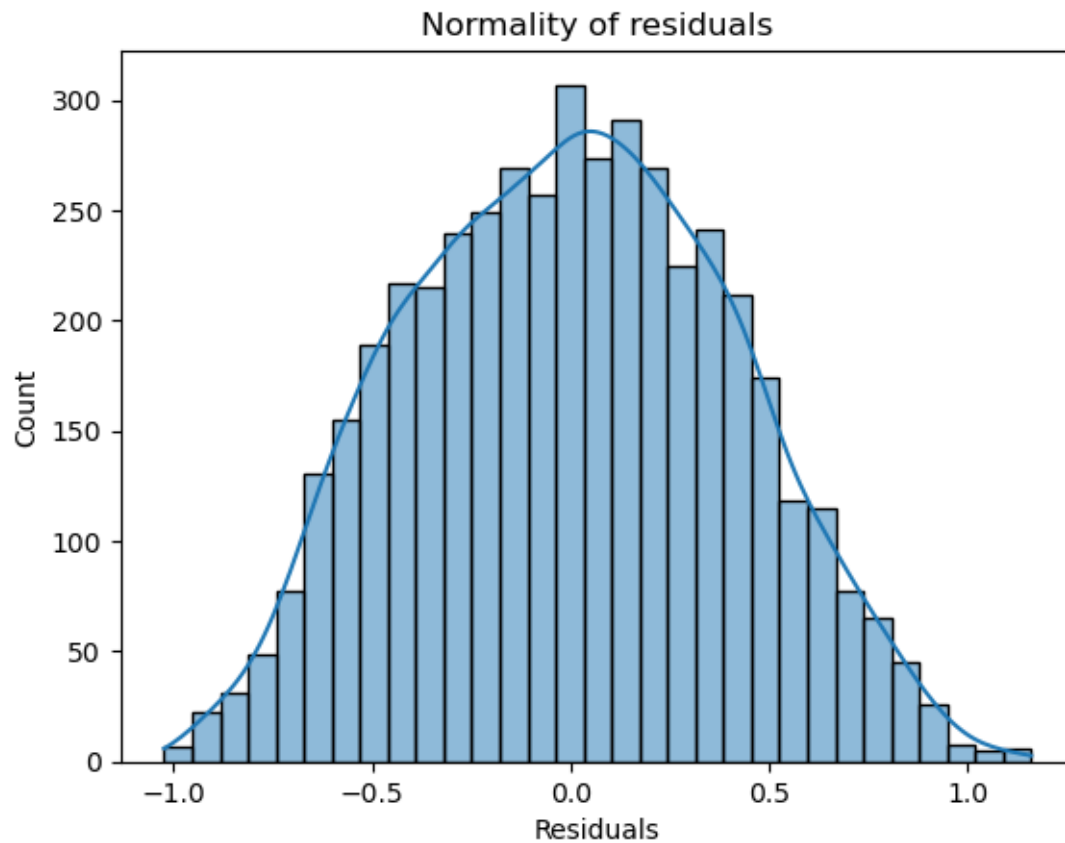
How to fix if this assumption is not followed?

- We can apply transformations like log, exponential, arcsinh, etc. as per our data.

```
[126]: import seaborn as sns
import matplotlib.pyplot as plt

# Plotting the histogram of residuals
sns.histplot(data=df_pred, x="Residuals", kde=True)
plt.title("Normality of residuals")
plt.show()
```


```
/opt/anaconda3/lib/python3.11/site-packages/seaborn/_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
```



- The histogram of residuals does have a bell shape.
- Let's check the Q-Q plot.

```
[129]: import pylab
import scipy.stats as stats

# Q-Q plot to check the normality of residuals
stats.probplot(df_pred["Residuals"], dist="norm", plot=pylab)
pylab.show()
```



- The residuals more or less follow a straight line except for the tails.
- Let's check the results of the Shapiro-Wilk test.

```
[132]: from scipy.stats import shapiro

# Perform the Shapiro-Wilk test on the residuals
shapiro_test = shapiro(df_pred["Residuals"])
print("Shapiro-Wilk Test:")
print("Test Statistic:", shapiro_test.statistic)
print("p-value:", shapiro_test.pvalue)
```

```
Shapiro-Wilk Test:
Test Statistic: 0.9934884905815125
p-value: 1.4591795924768364e-13
```

- Since $p\text{-value} < 0.05$, the residuals are not normal as per the Shapiro-Wilk test.
- Strictly speaking, the residuals are not normal.

- However, as an approximation, we can accept this distribution as close to being normal.
- So, the assumption is satisfied.

23 TEST FOR HOMOSCEDASTICITY

- **Homoscedascity:** If the variance of the residuals is symmetrically distributed across the regression line, then the data is said to be homoscedastic.
- **Heteroscedascity:** If the variance is unequal for the residuals across the regression line, then the data is said to be heteroscedastic.

Why the test?

- The presence of non-constant variance in the error terms results in heteroscedasticity. Generally, non-constant variance arises in presence of outliers.

How to check for homoscedasticity?

- The residual vs fitted values plot can be looked at to check for homoscedasticity. In the case of heteroscedasticity, the residuals can form an arrow shape or any other non-symmetrical shape.
- The goldfeldquandt test can also be used. If we get a p-value > 0.05 we can say that the residuals are homoscedastic. Otherwise, they are heteroscedastic.
- Null hypothesis: Residuals are homoscedastic
- Alternate hypothesis: Residuals have heteroscedasticity

How to fix if this assumption is not followed?

- Heteroscedasticity can be fixed by adding other important features or making transformations.

```
[139]: import statsmodels.stats.api as sms
from statsmodels.compat import lzip

# Define the names for the test statistics
names = ["F statistic", "p-value"]

# Perform the Goldfeld-Quandt test
test = sms.het_goldfeldquandt(y_train, X_train3)

# Print the results
results = dict(lzip(names, test))
```

```
print(results)
```

```
{'F statistic': 1.0484582413449557, 'p-value': 0.1304796784544968}
```

Since $p\text{-value} > 0.05$, we can say that the residuals are homoscedastic. So, this assumption is satisfied.

24 Predictions on test data

Now that we have checked all the assumptions of linear regression and they are satisfied, let's go ahead with prediction.

```
[144]: # Predictions on the test set
pred = olsmod2.predict(X_test3)

# Creating a DataFrame with actual and predicted values
df_pred_test = pd.DataFrame({"Actual": y_test, "Predicted": pred})

# Sampling 10 random rows from the DataFrame
print(df_pred_test.sample(10, random_state=1))
```

	Actual	Predicted
2023	2.039	2.262980
4543	3.236	3.784680
1914	2.314	2.323405
4856	2.457	2.176779
4405	2.322	2.331494
5755	4.344	3.928340
2341	2.858	2.279015
1637	3.976	3.879616
5730	3.728	3.588967
1385	2.692	2.420355

25 Final Model

Let's recreate the final model and print its summary to gain insights.

```
[147]: # Final training and testing sets
X_train_final = X_train3.copy()
X_test_final = X_test3.copy()

# Fit the final OLS model
olsmodel_final = sm.OLS(y_train, X_train_final).fit()
print(olsmodel_final.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          rating    R-squared:          0.718
```

```

Model:                OLS      Adj. R-squared:      0.717
Method:               Least Squares  F-statistic:      482.7
Date:                 Tue, 26 Nov 2024  Prob (F-statistic):      0.00
Time:                 23:21:21   Log-Likelihood:      -2340.4
No. Observations:     4566      AIC:              4731.
Df Residuals:         4541      BIC:              4891.
Df Model:              24
Covariance Type:      nonrobust

```

		coef	std err	t	P> t
[0.025 0.975]					

const		2.7875	0.033	84.367	0.000
2.723	2.852				
eps		0.0201	0.001	18.671	0.000
0.018	0.022				
duration		0.0123	0.000	26.100	0.000
0.011	0.013				
watched		0.0002	6.16e-06	31.750	0.000
0.000	0.000				
watching		0.0031	0.000	19.918	0.000
0.003	0.003				
years_running		-0.0762	0.009	-8.944	0.000
-0.093	-0.059				
mediaType_Movie		-0.3078	0.032	-9.550	0.000
-0.371	-0.245				
mediaType_Music Video		-0.2987	0.030	-9.971	0.000
-0.357	-0.240				
mediaType_OVA		-0.3016	0.030	-10.092	0.000
-0.360	-0.243				
mediaType_Other		-0.2607	0.035	-7.548	0.000
-0.328	-0.193				
mediaType_TV		-0.5598	0.033	-16.781	0.000
-0.625	-0.494				
mediaType_TV Special		-0.1816	0.039	-4.657	0.000
-0.258	-0.105				
mediaType_Web		-0.4164	0.031	-13.523	0.000
-0.477	-0.356				
contentWarn_Yes		-0.1786	0.020	-8.745	0.000
-0.219	-0.139				
studio_primary_J.C. Staff		-0.2017	0.041	-4.865	0.000
-0.283	-0.120				
studio_primary_MADHOUSE		-0.2532	0.043	-5.844	0.000
-0.338	-0.168				
studio_primary_OLM		-0.4121	0.052	-7.941	0.000
-0.514	-0.310				

studio_primary_Others	-0.2993	0.024	-12.445	0.000
-0.346 -0.252				
studio_primary_Studio Deen	-0.1657	0.048	-3.432	0.001
-0.260 -0.071				
studio_primary_Studio Pierrot	-0.2740	0.049	-5.591	0.000
-0.370 -0.178				
studio_primary_Sunrise	-0.1088	0.039	-2.815	0.005
-0.185 -0.033				
studio_primary_Toei Animation	-0.2325	0.034	-6.853	0.000
-0.299 -0.166				
genre_Adventure	-0.0783	0.032	-2.467	0.014
-0.140 -0.016				
genre_Comedy	-0.2271	0.052	-4.374	0.000
-0.329 -0.125				
genre_Drama	0.2881	0.040	7.192	0.000
0.210 0.367				

```
=====
Omnibus:                143.175    Durbin-Watson:                1.946
Prob(Omnibus):          0.000    Jarque-Bera (JB):            68.833
Skew:                   0.052    Prob(JB):                    1.13e-15
Kurtosis:               2.407    Cond. No.                     2.81e+04
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.81e+04. This might indicate that there are strong multicollinearity or other numerical problems.